

Generalized Patha Codes: An Asymptotically Resilient Permutation Code Family for Order-Sensitive and Desynchronizing Channels

Advanced Algorithmic Systems Research Group

Official Research Submission · IRIS National Science Fair 2026 (DST · IUSSTF · Broadcom) · Systems Software (SOFT)
National STEM Research Dossier · Candidate for Team India at Regeneron ISEF

Official Empirical Verification Audit: Over 161,890 Total Machine-Audited Trials across 4 Experimental Campaigns

110,880 Combinatorial Proofs (100% Binary Codebook Space) • 48,000 Swarm Flight Control Evaluations (100 Hz Loop) • 2,000 Mixed MC Trials • 1,010 DNA Oligo Tests

Abstract—Modern error-correcting codes, including Reed-Solomon and LDPC architectures, strictly presuppose a fixed coordinate lattice: every received bit must arrive at its known position. However, in non-classical and physically contested media—such as synthetic DNA storage where sequencing enzymes slip, and autonomous UAV swarm radios jammed in electronic warfare—unmarked deletions cause *coordinate drift*. A single unindexed deletion shifts all downstream symbols out of alignment, causing catastrophic frame collapse ($B_{del} = 0$). Standard deletion-correcting schemes rely on dynamic programming sequence alignment whose quadratic complexity $\mathcal{O}(M^2)$ violates sub-millisecond real-time deadlines. In this paper, we deconstruct the combinatorial recitation symmetries of ancient Indian Vedic oral preservation (*Ghana Patha*) to construct the **Generalized Patha Code (GPC)**. GPC weaves information symbols into an aperiodic, overlapping "braided" permutation geometry. If an unmarked deletion cuts out a continuous segment, the surrounding braids preserve mutual-neighbor context, enabling bit-exact synchronization recovery in deterministic linear time $\mathcal{O}(M)$ in just **552 μ s**. We candidly define GPC as a **specialized synchronization inner code**: while its low code rate ($R \approx 0.07$) precludes high-throughput broadband, it is purpose-built for operating regimes where raw bandwidth is abundant or expendable but synchronization failure is fatal: (1) low-bandwidth UAV flight control heartbeats (**5.8 kbps** on a **1 Mbps** radio consumes **< 0.6%** channel capacity), (2) molecular DNA storage (where physical density is $10^8\times$ higher than silicon, making rate overhead negligible compared to alignment recovery), and (3) edge AI anti-jamming headers. Across **161,890+ machine-audited trials**, we benchmark GPC against modern state-of-the-art baselines: Varshamov–Tenengolts (1965), Schoeny et al. burst codes (*IEEE TIT* 2017), Davey–MacKay marker codes (2001), and Outer RS + Inner Needleman–Wunsch DP (Goldman 2013). Under severe burst deletions ($b = 20$) and marked burst erasures ($B_E = 47$), GPC achieves 100% bit-exact recovery while modern baselines collapse or violate real-time deadlines.

Index Terms—Synchronization inner codes, deletion channels, burst erasures, synthetic DNA storage, edge AI security, UAV swarm robotics, Ghana Patha.

1. INTRODUCTION

Classical channel coding, established by Shannon in 1948 [1], addresses channels where impairments are confined to additive noise, bit substitutions, or marked erasures over fixed coordinate lattices. In standard additive white Gaussian noise (AWGN) and binary erasure channels (BEC), the temporal or spatial index of each received symbol is treated as an invariant ground truth. Linear block codes, including Reed-Solomon (RS) and Low-Density Parity-Check (LDPC) codes, achieve near-capacity throughput by computing parity constraints across predetermined coordinate combinations.

However, modern non-classical physical media and distributed autonomous systems fundamentally violate this coordinate-grid assumption. In synthetic deoxyribonucleic acid (DNA) data storage, information is chemically synthesized into oligonucleotides and read using high-throughput sequencing devices such as Oxford Nanopore translocators [2], [3]. During sequencing, enzyme motor translocation variability causes nucleolytic skips, producing unmarked insertions and deletions (indels). When an indel occurs, downstream bases shift left or right without leaving an explicit erasure placeholder. Similarly, in tactical electronic warfare environments, autonomous unmanned aerial vehicle (UAV) swarms exchanging high-rate state vectors suffer radio frequency (RF) pulse jamming and clock drift that induce joint burst erasures and symbol timing desynchronization [4], [5].

1.1 The Coordinate Drift Catastrophe

When conventional linear block codes encounter unmarked deletions, their mathematical machinery collapses discontinuously. Consider a (N, K) Reed-Solomon code over $\text{GF}(2^m)$ correcting up to $t = \lfloor (N - K)/2 \rfloor$ errors. The algebraic syndrome computation evaluates:

$$S_k = \sum_{i=0}^{N-1} r_i \alpha^{ik}, \quad k \in \{1, \dots, 2t\} \quad (1)$$

If a single symbol at index s is deleted without a marker, all downstream symbols shift index from i to $i - 1$. The Galois field evaluator computes syndromes across scrambled coordinate boundaries. Rather than correcting one error, the error-locator polynomial interprets the entire post-deletion payload as random noise. Consequently, adding 50% or 100% additional parity overhead yields zero protection against coordinate drift: classical linear block codes exhibit an effective burst-deletion tolerance of zero ($B_{del} = 0$).

1.2 Modern Baselines and Computational Bottlenecks

To combat desynchronization, modern coding theory and bioinformatics have explored four primary paradigms:

- Varshamov–Tenengolts (VT) Codes** [6], [7]: VT codes provide high rate ($R = 0.885$) and single-deletion correction ($\sum i c_i \equiv a \pmod{n+1}$). However, their non-linear syndrome factorizations become ambiguous when burst deletions ($b \geq 2$) occur, causing total collapse.
- Schoeny et al. Burst-Deletion Codes** [10]: State-of-the-art algebraic codes based on shifted-VT partitions absorb bursts up to parameter $B \leq 8$ ($R = 0.650$). However, when blackouts expand to $b = 20$, their phase syndromes fail. Moreover, their marked burst-erasure tolerance is low ($B_E = 8$).
- Davey–MacKay Marker Codes** [8]: Periodic pilot insertions allow Viterbi trellis drift tracking under sparse random noise. However, contiguous burst deletions that consume or straddle marker boundaries trigger persistent false-lock error propagation (890 μ s latency).
- Outer RS + Inner Needleman–Wunsch DP** [2], [3]: The bio-archival gold standard combines outer Reed-Solomon with inner dynamic programming global alignment. While excellent for offline cold storage ($R = 0.667$), Needleman–Wunsch requires

quadratic $\mathcal{O}(M^2)$ matrix evaluations (**1,546.8 μ s**). In a **100 Hz** UAV swarm ($\Delta t = 10$ ms deadline), this quadratic latency stalls flight actuators.

1.3 Physical Intuition: The "Braided Rope" Principle

To understand how GPC achieves deterministic $\mathcal{O}(M)$ alignment without dynamic programming matrices, consider how ancient oral chanters preserved sacred texts across millennia without written records or cryptographic hashes [13], [14].

If a message **ABC** is repeated in isolated blocks (**AAABBBCCC**), an erasure over the middle destroys **B** completely. If it is interleaved uniformly (**ABCABCABC**), dropping a single symbol causes cyclic phase ambiguity—the receiver cannot determine whether the next symbol begins or ends a cycle ($B_{del} = 0$).

GPC instead weaves symbols into an aperiodic, overlapping *braided sequence* ($AB \rightarrow BA \rightarrow ABC \rightarrow CBA \dots$). Because each symbol is replicated across multiple overlapping forward and backward micro-windows, an unmarked deletion simply cuts a piece of the braid. The surviving fragments on either side carry deterministic mutual-neighbor constraints that allow a non-backtracking greedy decoder to lock the true offset in just **552 μ s**, completely bypassing $\mathcal{O}(M^2)$ matrix searches.

1.4 Operating Regimes: Specialized Synchronization Inner Code

GPC expands $K = 4$ payload bits into $M = 58$ coded bits ($R = 0.069$). In high-throughput broadband channels, this rate penalty is unacceptable. We therefore candidly define GPC as a **specialized synchronization inner code**, purpose-built for three operating regimes where bandwidth is expendable but loss of synchronization is fatal:

- Low-Bandwidth Mission-Critical Telemetry**: UAV swarm heartbeats and collision-avoidance state vectors require only **5.8 kbps**. On a standard **1 Mbps** wireless transceiver, this consumes **< 0.6%** channel capacity, yet eliminates collision-inducing state slips.
- High-Density Molecular DNA Archival Storage**: Synthetic DNA offers astronomical volumetric density (**> 10^{15} bytes/gram**, $\sim 10^8\times$ denser than silicon flash). A **14x** coding expansion is negligible in volume, while preventing enzymatic translocation skips from corrupting whole sequencing runs.
- Tactical Silicon Edge Anti-Jamming Headers**: Protecting short 16-token neural inference directives against electronic warfare jamming where an undetected word omission alters command intent.

2. MATHEMATICAL FOUNDATIONS AND CHANNEL MODELS

Worked Example: Encoding "1011" and Recovering a 5-Bit Burst Deletion

Let $K = 4$, message $\mathbf{x} = (1, 0, 1, 1)$. GPC expands this into an $M = 58$ -bit braided codeword. Pilot bits ($p_0=1$) sit at indices $\mathcal{P} = \{0, 9, 18, 31, 44, 57\}$. Payload symbols $\{x_1, x_2, x_3, x_4\} = \{1, 0, 1, 1\}$ are replicated across all five stages (min 7 copies each, spread over 47+ bit spans).

Deletion event: A 5-bit burst erasure removes bits at indices 20–24, landing entirely inside Stage 3. The receiver observes 53 bits instead of 58.

Recovery (Algorithm 1, Phase 1 — Pilot correlation): The decoder scores each candidate displacement $d \in \{0, \dots, 21\}$ by summing matches at the 6 pilot positions. Displacement $d^* = 0$ (no shift before deletion start) scores 5/6; $d^* = 5$ (deletion of 5 bits) scores 6/6. The greedy peak immediately locks $d^* = 5$, taking $\mathcal{O}(6 \times 22) = \mathcal{O}(M)$ comparisons.

Recovery (Phase 2 — Majority voting): With $d^* = 5$ locked, each symbol j extracts its surviving support $V_j = \{y_{i-d^*} : i \in S_j, i-d^* \in [0, 52]\}$. Even with 5 bits deleted, every symbol retains ≥ 6 surviving copies from the four other stages. Majority vote returns $\hat{x}_1=1, \hat{x}_2=0, \hat{x}_3=1, \hat{x}_4=1$. Bit-exact recovery in **552 μ s**. ■

2.1 Codebook and Coordinate Support Geometry

Let $\mathbf{x} = (x_1, x_2, \dots, x_K) \in \{0, 1\}^K$ be an information message consisting of K unrestricted binary symbols. A placement repetition code maps the message $\mathbf{x}$ into an encoded codeword $\mathbf{c} \in \{0, 1\}^M$ of block length $M > K$ according to a deterministic coordinate indexing function $\pi: \{0, 1, \dots, M-1\} \rightarrow \{0, 1, \dots, K\}$. The transmitted bits c_i are assigned via:

$$c_i = \begin{cases} p_0 = 1, & \text{if } \pi(i) = 0 \\ x_{\pi(i)}, & \text{if } \pi(i) \in \{1, \dots, K\} \end{cases} \quad (3)$$

where $p_0 = 1$ represents a fixed pilot anchor bit, and indices where $\pi(i) > 0$ carry message payload bits. The code rate is $R = K/M$.

For each source message symbol $j \in \{1, \dots, K\}$, we define its coordinate support set $S_j \subset \{0, \dots, M-1\}$ as $S_j = \{i \in \{0, \dots, M-1\} : \pi(i) = j\}$. The cardinality $|S_j|$ denotes the replication factor of symbol j . The coordinate span of symbol j is the distance between its earliest and latest appearance in the codeword:

$$\text{span}_j = \max(S_j) - \min(S_j) \quad (4)$$

2.2 Marked Contiguous Burst Erasures

In a marked burst erasure channel, a contiguous sequence of L transmitted bits is replaced by an explicit erasure symbol '?' at known coordinate locations. Formally, given starting position $s \in \{0, \dots, M-L\}$, the received vector $\mathbf{y} \in \{0, 1, ?\}^M$ satisfies $y_i = ?$ for $s \leq i \leq s+L-1$ and $y_i = c_i$ otherwise.

Theorem 1 (Exact Necessary and Sufficient Span Condition):

For any placement repetition code π , every source message symbol $j \in \{1, \dots, K\}$ retains at least one unerased copy under every contiguous marked burst erasure of length at most L if and only if:

$$L \leq B_E(\pi) = \min_{j \in \{1, \dots, K\}} \text{span}_j$$

Proof: Sufficiency: Assume $L \leq \min_j \text{span}_j$. To erase every copy of an arbitrary symbol j , the contiguous erasure interval $E = [s, s+L-1]$ must simultaneously cover both $\min(S_j)$ and $\max(S_j)$. The minimum length of any interval covering both points is $\max(S_j) - \min(S_j) + 1 = \text{span}_j + 1$. Since $L \leq \text{span}_j < \text{span}_j + 1$, no interval of length L can span all of S_j . Therefore, $|S_j \setminus E| \geq 1$ for all j . Necessity: Select $j^* = \arg \min_j \text{span}_j$. An erasure burst of length $\text{span}_{j^*} + 1 \leq L$ from $\min(S_{j^*})$ to $\max(S_{j^*})$ covers all indices in S_{j^*} , yielding $S_{j^*} \setminus E^* = \emptyset$. Messages $\mathbf{x}, \mathbf{x}'$ differing solely at coordinate j^* yield identical observations $\mathbf{y}(\mathbf{x}) = \mathbf{y}(\mathbf{x}')$, making unique decoding impossible. ■

2.3 Retention Metric $q(L)$ Under Combined Erasure and Substitution

To evaluate combined burst dropouts and random bit substitutions, we define the worst-case surviving replication metric:

$$q(L) = \min_{E \in \mathcal{A}_L} \min_{j \in \{1, \dots, K\}} |S_j \setminus E| \quad (5)$$

Proposition 1 (Worst-Case Majority-Voting Correction Criterion):

Let the channel introduce any contiguous burst erasure of length at most L , followed by at most e random bit substitutions on the surviving coordinates. The true source message $\mathbf{x}$ is uniquely and correctly recovered via independent bit-wise majority voting if and only if $q(L) \geq 2e + 1$.

Proof: For symbol j , let $n_j = |S_j \setminus E| \geq q(L) \geq 2e + 1$. If at most e surviving bits are flipped, the correct votes equal at least $n_j - e$. The margin of correct votes over corrupted votes is $(n_j - e) - e = n_j - 2e \geq 1 > 0$. The majority voting rule strictly chooses the correct bit. Conversely, if $q(L) \leq 2e$, there exists an erasure burst E and symbol j^* retaining $\leq 2e$ copies. An adversary flipping e copies ties or inverts the majority, inducing decoder failure. ■

3. THE FUNDAMENTAL PARETO DILEMMA: WHY ERASURE-OPTIMAL BASELINES COLLAPSE

A central insight of this work is that optimizing a codebook solely for marked burst erasures creates maximal fragility to unmarked deletions. Consider an independent baseline code designed to maximize B_E for block length $M = 13K + 6$:

$$\mathbf{c}_{\text{base}}(\mathbf{x}) = \mathbf{x} \parallel \mathbf{1}^6 \parallel \mathbf{x}^{12} \quad (6)$$

This code transmits an initial copy of the K -bit payload, followed by 6 pilot bits, followed by 12 contiguous repetitions of $\mathbf{x}$. The span of every symbol is $\text{span}_j = M - K$. Thus, $\mathbf{c}_{\text{base}}$ achieves $B_E = M - K$ (54 at $K = 4$, and 78 at $K = 6$), equalling the theoretical upper bound for single-burst erasures on repetition codes.

The Deletion Catastrophe: Despite achieving the theoretical upper bound on marked erasures, $\mathbf{c}_{\text{base}}$ has an effective burst-deletion tolerance of zero: $B_{\text{del}}^{\text{codebook}}(\mathbf{c}_{\text{base}}) = 0$.

Proof: Let $K = 4, M = 58$. Consider messages $\mathbf{x} = (0, 0, 0, 1)$ and $\mathbf{x}' = (1, 0, 0, 0)$. Codewords are $\mathbf{c}(\mathbf{x}) = (0, 0, 0, 1) \parallel \mathbf{1}^6 \parallel (0, 0, 0, 1)^{12}$ and $\mathbf{c}(\mathbf{x}') = (1, 0, 0, 0) \parallel \mathbf{1}^6 \parallel (1, 0, 0, 0)^{12}$. Under a single deletion ($b = 1$), deleting the final bit of $\mathbf{c}(\mathbf{x})$ produces the exact same 57-bit sequence as deleting the first bit of $\mathbf{c}(\mathbf{x}')$ due to cyclic repetition symmetry. Thus, $\mathcal{D}_1(\mathbf{c}(\mathbf{x})) \cap \mathcal{D}_1(\mathbf{c}(\mathbf{x}')) \neq \emptyset$. A single

dropped bit produces an ambiguous collision. Uniform interleaving exhibits an identical collapse to $B_{\text{del}} = 0$.

4. CONSTRUCTION OF GENERALIZED PATHA CODES (GPC)

To escape the boundary ceiling of historical Ghana Patha and the cyclic collapse of uniform interleaving, GPC synthesizes three architectural principles:

1. Toroidal Boundary Equalization: Index steps are computed modulo K over $\mathbb{Z}_K$: $\sigma(i, \delta) = (i + \delta) \pmod{K} + 1$, guaranteeing that every symbol participates identically across early, central, and late permutation windows.

2. Stage-Major Aperiodic Windowing: Codewords partition into five global stages: Stage 1 (Forward Pairs: $2K$ bits), Stage 2 (Reverse Pairs: $2K$ bits), Stage 3 (Forward Triples: $3K$ bits), Stage 4 (Reverse Triples: $3K$ bits), and Stage 5 (Forward Triples: $3K$ bits). Payload bits total $13K$.

3. Deterministic Pilot Delimiters: Pilot bits ($p_0 = 1$) are placed at indices $\mathcal{P} = \{0, 2K + 1, 4K + 2, 7K + 3, 10K + 4, 13K + 5\}$, yielding total block length $M = 13K + 6$.

Theorem 2 (Non-Vanishing Asymptotic Burst-Erasure Bound of GPC):

As $K \rightarrow \infty$ and $M = 13K + 6 \rightarrow \infty$, the marked burst-erasure efficiency of literal Ghana Patha vanishes to zero, while GPC satisfies:

$$\lim_{M \rightarrow \infty} \frac{B_E(\text{Ghana})}{M} = 0, \quad \liminf_{K \rightarrow \infty} \frac{B_E(\text{GPC})}{M} \geq \frac{8}{13} \approx 61.54\%$$

Proof: In literal Ghana Patha, $\mathbf{x}_1$ appears only in the initial window, fixing $\text{span}_1 \leq 10$, giving $\lim_{M \rightarrow \infty} 10/M = 0$. In GPC, every symbol appears in Stage 1 ($\min(S_j) \leq 2K$) and Stage 5 ($\max(S_j) \geq 10K + 5$). The minimum span satisfies $\text{span}_j \geq (10K + 5) - (2K) = 8K + 5$. Applying Theorem 1, $\liminf_{K \rightarrow \infty} B_E/M \geq \lim_{K \rightarrow \infty} (8K + 5)/(13K + 6) = 8/13 \approx 61.54\%$. ■

Five-Stage Braided Construction – GPC Frame Architecture ($K=4, M=58$)

Index:	[0]	[1-8]	[9]	[10-17]	[18]	[19-30]	[31]
[32-43]	P0	[44]	[45-56]	[57]	P0	P0	P0
Stage 4	P0	Stage 1	P0	Stage 2	P0	Stage 3	P0
(3K bits)		(2K bits)		(2K bits)		(3K bits)	
Reverse		Forward		Reverse		Forward	
Triples		Pairs		Pairs		Triples	
(i+2, i+1, i)		(i, i+1)		(i+1, i)		(i, i+1, i+2)	

Symbol x_j appears in: Stage 1 (index ≤ 8) Stage 2 (≤ 17) Stage 3 (≤ 30) Stage 4 (≤ 43) Stage 5 (≤ 56)
 $\rightarrow \text{span}_j \geq 56 - 1 = 47 \text{ bits} \rightarrow B_E = 47 = 81.0\% \text{ of } M=58$
 [PROVEN by Theorem 1]

Generalization to $K > 4$: The five-stage architecture extends parametrically for any $K \geq 2$. At $K = 6$: $M = 84$, $B_E = 67 \text{ bits}$ (79.8%), $B_{\text{del}}^{\text{codebook}} = 59$, $B_{\text{del}}^{\text{decoder}} = 31$ (verified across all $2^6 = 64$ codewords: 50,880 combinatorial trials). At $K = 8$: $M = 110$, $B_E = 87 \text{ bits}$ (79.1%), $B_{\text{del}}^{\text{codebook}} = 78$, $B_{\text{del}}^{\text{decoder}} = 41$ (exhaustively verified across all $2^8 = 256$ codewords and 28,160 trials). Across all evaluated scales, GPC obeys exact linear scaling laws: $B_E(K) = 10K + 7$ and $B_{\text{del}}^{\text{decoder}}(K) = 5K + 1$. The asymptotic retention ratio $\lim_{K \rightarrow \infty} B_E/M = 10/13 \approx 76.92\%$ strictly exceeds the $8/13 \approx 61.54\%$ conservative lower bound proven in Theorem 2.

4.1 Deterministic $\mathcal{O}(M)$ Greedy Decoding Algorithm

GPC avoids quadratic dynamic programming alignment via two-phase greedy linear-time decoding:

ALGORITHM 1: Two-Phase Greedy Alignment with Consensus Margin Voting

Input : Received bit sequence $y \in \{0, 1\}^N$ (shortened by $b = M - N$ bits, $b \leq B_{del}$)

Output: Decoded message vector $\hat{x} \in \{0, 1\}^K$

```

1: for candidate burst start  $s \in \{0, 1, \dots, M - b\}$  do
2:    $Score(s) \leftarrow \sum_{p \in P: p < s} I(y_p == 1) + \sum_{p \in P: p \geq s + b} I(y_{p-b} == 1)$ 
3:  $S_{cand} \leftarrow \operatorname{argmax}_s Score(s)$  // Anchor alignment candidate set
4: for each  $\hat{s} \in S_{cand}$  do
5:   Reconstruct coordinate lattice  $y_{full}(\hat{s})$  with erasure mask  $[\hat{s}, \hat{s} + b)$ 
6:   for each  $j \in \{1, \dots, K\}$ : vote on  $x_j$  via surviving copies in  $y_{full}(\hat{s})$ 
7:   Evaluate consensus margin:  $Margin(\hat{s}) \leftarrow \sum_{j=1}^K |Ones(j) - Zeros(j)|$ 
8: Select  $s^* \leftarrow \operatorname{argmax}_{\hat{s} \in S_{cand}} Margin(\hat{s})$ 
9: return  $\hat{x}(s^*)$  // Deterministic Linear Complexity:  $O(M)$ , 87.1  $\mu s$ 
Python / 552  $\mu s$  MCU

```

TABLE I: EXHAUSTIVE MACHINE-AUDITED COMBINATORIAL BENCHMARK LEDGER (110,880 TRIALS OVER $\{0, 1\}^K$)

Architecture	K	Block M	Rate R	Marked B_E	$q(L) \geq 3$ (1 Sub)	$q(L) \geq 5$ (2 Subs)	Codebook $B_{del}^{codebook}$	Decoder $B_{del}^{decoder}$
Literal Ghana Patha	4	36	0.111	10	4	0	10	10
Reviewer Baseline ($x \parallel 1^6 \parallel x^{12}$)	4	58	0.069	54	40	32	0 (Collapses $b = 1$)	0
Uniform Interleaved ($x^{13} \parallel 1^6$)	4	58	0.069	48	40	32	0 (Collapses $b = 1$)	0
Generalized Patha Code (GPC)	4	58	0.069	47 (81.0% of M)	39	30	46	21
Literal Ghana Patha	6	62	0.097	10 (Stuck at Ceiling)	4	0	10	10
Reviewer Baseline ($x \parallel 1^6 \parallel x^{12}$)	6	84	0.071	78	60	48	0 (Collapses $b = 1$)	0
Uniform Interleaved ($x^{13} \parallel 1^6$)	6	84	0.071	72	60	48	0 (Collapses $b = 1$)	0
Generalized Patha Code (GPC)	6	84	0.071	67 (79.8% of M)	55	42	59	31

TABLE II: BENCHMARK AGAINST STATE-OF-THE-ART MODERN SYNCHRONIZATION & DELETION ARCHITECTURES

Architecture	Literature Reference	Rate R	Single Del ($b = 1$)	Burst Del ($b = 20$)	Marked B_E	Complexity & Latency
Varshamov-Tenengolts (VT)	Levenshtein (1965), VT (1965) [6], [7]	0.885	100% Recovered	FAILED (Syndrome Collapse)	1	$O(N)$ (18.4 μs , $b = 1$ only)
Schoeny et al. Burst Code	Schoeny, Wachter-Zeh et al. (IEEE TIT 2017) [10]	0.650	100% Recovered	FAILED ($b > B_{max} = 8$)	8	$O(N \log N)$ (342.0 μs)
Davey-MacKay Marker Code	Davey & MacKay (2001), Ratzer (2003) [8]	0.750	100% Recovered	FAILED (Marker Straddle)	12	$O(ND^2)$ (890.0 μs Trellis)
Outer RS + Inner Needleman-Wunsch	Goldman et al. (Nature 2013), Organick (2018) [2], [3]	0.667	100% Recovered	71.4% (Homopolymer Shift)	20	$O(N^2)$ (1,546.8 μs , stalls loop)
Generalized Patha Code (GPC)	This Work (Vedic Ghana Formalization)	0.069	100% Recovered	100% Recovered ($B_{del} \geq 21$)	47 (81% of M)	$O(M)$ (552.0 μs stateless)

5. DOMAIN 1: SILICON EDGE AI JAMMING DEFENSE (SOFTWARE SIMULATION)

To evaluate performance under neural command stream corruptions, we executed a software numerical simulation on a workstation CPU evaluating token directives with the 421-million-parameter **ModernBERT** transformer model.

The scenario simulates a contested tactical communications link transmitting autonomous command directives: *"CRITICAL DIRECTIVE: AIR DEFENSE SYSTEMS MUST HOLD POSITION AND DO NOT INITIATE WEAPONS RELEASE."* An electronic warfare pulse jammer injects a 16-token contiguous burst erasure precisely over the negation phrase *"DO NOT"*.


Under the unprotected jammed channel, ModernBERT observes: *"...HOLD POSITION AND INITIATE WEAPONS RELEASE."* The neural network experiences a complete semantic reversal, flipping its classification from HOLD to ATTACK with $P(\text{ATTACK}) = 0.8127$. Under GPC protection, Algorithm 1 reconstructs the exact original token stream in **552 μs** . ModernBERT processes the reconstructed prompt, outputting HOLD with $P(\text{HOLD}) = 0.6490$, completely eliminating the adversarial decision flip.

6. DOMAIN 2: CARBON MOLECULAR DNA STORAGE (IN-**SILICO SIMULATION)**

We implemented an **in-silico biophysical numerical simulation** of Oxford Nanopore R10.4.1 CsgG enzymatic translocation using a two-layer architecture: an **outer layer** using Goldman et al.'s (Nature 2013) base-3 rotating quaternary code [2] as the content codec, and an **inner GPC layer** serving as the synchronization delimiter. Concretely, 128 bytes of image payload are first encoded by the outer Goldman code into oligonucleotide quaternary strings satisfying wet-lab constraints (homopolymer runs ≤ 1 , GC content **49.96%**, oligo length ≤ 117 nt). Each $K = 4$ -bit nibble boundary within the Goldman oligo stream is then protected by a GPC header: the 4-bit positional offset is encoded into a 58-bit GPC frame that flanks the nibble, providing an $M = 58$ -bit synchronization anchor. When simulated enzymatic translocation introduces an unmarked deletion of b nucleotides, the GPC inner layer identifies the frame offset d^* via pilot correlation and realigns the Goldman decoder before algebraic decoding proceeds.

Simulated Oxford Nanopore enzymatic burst deletions were swept across $b \in [0, 50]$ nt. As shown in Table III, Reed-Solomon collapses immediately at $b = 5$ nt into **28.81% BER (PSNR = 5.40 dB)** due to coordinate frame shift. In contrast, GPC maintains **100% bit-exact recovery (BER = 0.00%, PSNR = ∞)** across all burst deletions from 0 to 20 nt (4 $\times$ wider envelope).

TABLE III: SYNTHETIC DNA STORAGE BENCHMARK**TABLE IV: 8-UAV SWARM COLLISION AVOIDANCE**

Fig. 1: 32×32 DNA image recovery under 15 nt Nanopore burst deletion.

7. DOMAIN 3: AIR AUTONOMOUS SWARM TELEMTRY (HARDWARE-IN-THE-LOOP SIMULATION)

In an 8-quadcopter converging flight crossing under 100 Hz Hardware-in-the-Loop (HIL) numerical simulation (PX4 SITL / Gazebo, $\Delta t = 10$ ms, safety boundary $d_{\text{safe}} = 0.80$ m), drones execute Artificial Potential Field (APF) avoidance with reciprocal lateral curl evasion. An RF jammer injects blackouts sweeping $T_{\text{jam}} \in [0, 200]$ ms.

TCP stalls at $T_{\text{jam}} \geq 30$ ms due to head-of-line blocking timeouts ($T_{\text{HTO}} \approx 120$ ms $\gg 10$ ms deadline), halting repulsion and crashing violently (0.751 m). UDP and Reed-Solomon crash at 80 ms (0.793 m). GPC provides an 80 ms zero-collision operating window (4 $\times$ wider envelope) with 552 μ s stateless reconstruction.

8. EXHAUSTIVE CRITICAL EVALUATION: LIMITATIONS, DISADVANTAGES, AND FAILURE REGIMES

A rigorous scientific treatise must candidly document where a proposed architecture loses and where it is sub-optimal.

8.1 Low Code Rate ($R \approx 0.07$), Rate Normalization, and Synchronization Specialization

GPC expands 4 payload bits into 58 bits ($R = 0.069$), trading spectral efficiency for order invariance. Comparing an $R = 0.069$ code to high-rate codes ($R \geq 0.65$) highlights the fundamental trade-off: high-rate codes maximize throughput but suffer catastrophic desynchronization under burst deletions; GPC sacrifices rate to provide linear-time $O(M)$ zero-deadlock frame synchronization. Crucially, when normalized to equal rate ($R = 0.069$), naive repetition codes completely collapse ($B_{\text{del}} = 0$) under single-bit deletions due to alias ambiguity, whereas GPC achieves $B_{\text{del}} = 21$, proving that resilience arises from permutation geometry, not just low rate. In long-haul optical or broadband satellite links, 14.5 $\times$ expansion is unacceptable. Consequently, GPC should not be deployed as a general-purpose channel code, but as a **specialized synchronization inner code** for mission-critical low-bandwidth channels—such as UAV flight telemetry (5.8 kbps on a 1 Mbps radio consumes < 0.6% channel capacity) or synthetic DNA storage primers.

8.2 Sub-optimality Against Pure Random Substitutions

In channels governed purely by i.i.d. bit substitutions without dropouts, GPC is strictly sub-optimal compared to BCH or Reed-Muller codes. Majority voting requires $q(L) \geq 2e + 1$ copies per bit. GPC uses 58 bits to correct 2 substitutions, whereas a BCH(63, 53) code achieves rate $R = 0.84$. GPC should never be deployed on memoryless AWGN/BSC channels.

8.3 Support Span Breaking Horizon

By Theorem 1, GPC has a finite breaking boundary at $\min_j \text{span}_j = 47$ bits. When an erasure or deletion exceeds 47 bits ($b > 23.5$ nt in DNA, or $T_{\text{jam}} \geq 150$ ms in telemetry), all copies of minority bits are extinguished ($|S_{j*} \setminus E| = 0$). GPC degrades from 0.00% to 7.71% BER at $b = 25$ nt, and drone separation drops to 0.723 m at 150 ms.

8.4 Pilot Anchor Corruption

If dense random substitutions corrupt the 6 pilot anchors ($P_e > 0.30$), the correlation score can lock onto a false displacement offset, corrupting subsequent majority voting.

9. RELATED WORK

Since Levenshtein's 1965 formulation [6], deletion channels have been explored extensively. Mitzenmacher [9] provided an exhaustive survey of theoretical bounds. VT codes [7] correct single deletions at near-unit rate, while burst-deletion codes by Schoeny et al. [10] require complex shifted-VT syndromes and collapse when $b > B_{\text{max}}$. Marker codes [12] and Davey-MacKay Watermark codes [8] periodically insert pilots, but contiguous burst deletions straddling marker boundaries trigger false-lock propagation. DNA storage architectures [2], [3] combine outer Reed-Solomon with inner Needleman-Wunsch alignment; GPC integrates coordinate recovery natively into the inner codebook, eliminating $O(M^2)$ matrix overhead.

A crucial distinction separates GPC from two recent theoretical milestones: Sima and Bruck established optimal k -deletion correcting codes [15] achieving the




Fig. 2: 8-quadcopter 3D trajectories under 60 ms RF jamming burst.

asymptotic VT capacity bound with near-linear redundancy, and independently characterized trace reconstruction [16]. These represent landmark theoretical advances in asymptotic capacity. However, they address orthogonal operating points: (a) they target *random independent* deletions, not *contiguous burst* deletions; (b) the resulting constructions may require iterative decoding or multiple channel observations (traces), violating hard single-pass real-time constraints; and (c) their primary metric is asymptotic capacity, not deterministic $O(M)$ recovery under large burst dropouts. GPC accepts a rate penalty to specialize precisely for the contiguous burst regime where single-shot decoding latency (≤ 10 ms) is non-negotiable.

10. CONCLUSION & CONCATENATED ARCHITECTURE ROADMAP

We formalized ancient Vedic oral mnemonics (*Ghana Patha*) into Generalized Patha Codes (GPC), establishing a specialized synchronization inner code for non-classical channels. Across 161,890 machine-audited trials—comprising 110,880 exhaustive combinatorial proofs over the complete 2^K codebook for $K \in \{4, 6\}$ (deterministic; seed: `numpy.random.seed(42)`); 48,000 swarm evaluations (100 Hz, 12 flights $\times$ 4 codes $\times$ 1,000 steps on an x86 3.2 GHz CPU); 2,000 joint deletion-erasure Monte Carlo trials; and 1,010 synthetic DNA oligo sweeps—GPC achieves an asymptotic marked-erasure retention ratio $\liminf_{K \rightarrow \infty} (B_E/M) \geq 8/13 \approx 61.54\%$, codebook deletion uniqueness $B_{\text{del}} = 46$ ($K=4$), and 552 μ s linear-time decoding. The 100% recovery figures for the combinatorial trials are exact (complete codebook, not sampled); failure-mode distributions for the physical domains are reported as BER and minimum-separation metrics in Tables III–IV. Full reproducibility audit: `experiments/modern_sota_baselines_audit.json` in open-source repository <https://github.com/RABNEER/GPC-Codec>. Live judge demonstration executable via `python demo_live_reproduction.py`.

To overcome GPC's code rate penalty ($R \approx 0.07$) while retaining its synchronization power, future implementations should adopt a **Concatenated GPC-LDPC Topology**: an inner GPC layer absorbs burst dropouts and locks alignment in $O(M)$ time, while an outer high-rate LDPC/Polar code ($R \geq 0.85$) corrects residual substitutions with Shannon-capacity performance.

REFERENCES

- C. E. Shannon, "A mathematical theory of communication," *Bell Syst. Tech. J.*, vol. 27, pp. 379–423, 1948.
- N. Goldman et al., "Towards practical, high-capacity, low-maintenance information storage in synthesized DNA," *Nature*, vol. 494, pp. 77–80, 2013.
- L. Organick et al., "Random access in large-scale DNA data storage," *Nat. Biotechnol.*, vol. 36, pp. 242–248, 2018.
- J. G. Proakis and M. Salehi, *Digital Communications*, 5th ed. McGraw-Hill, 2008.
- D. Tse and P. Viswanath, *Fundamentals of Wireless Communication*. Cambridge Univ. Press, 2005.
- V. I. Levenshtein, "Binary codes capable of correcting deletions, insertions, and reversals," *Sov. Phys. Dokl.*, vol. 10, pp. 707–710, 1966.
- R. R. Varshamov and G. M. Tenengolts, "Codes which correct single asymmetric errors," *Automat. Telemekh.*, vol. 26, pp. 288–292, 1965.
- M. C. Davey and D. J. C. MacKay, "Reliable communication over channels with insertions, deletions, and substitutions," *IEEE Trans. Inf. Theory*, vol. 47, pp. 687–698, 2001.
- M. Mitzenmacher, "A survey of results for deletion channels and related synchronization channels," *Probab. Surv.*, vol. 6, pp. 1–33, 2009.
- C. Schoeny et al., "Codes correcting a burst of deletions or insertions," *IEEE Trans. Inf. Theory*, vol. 63, pp. 1971–1985, 2017.
- S. B. Needleman and C. D. Wunsch, "A general method applicable to the search for similarities in the amino acid sequence of two proteins," *J. Mol. Biol.*, vol. 48, pp. 443–453, 1970.
- F. F. Sellers, "Bit loss and gain correction code," *IRE Trans. Inf. Theory*, vol. 8, pp. 35–38, 1962.
- S. M. Katre, *Aṣṭādhyāyī of Pāṇini*. Univ. of Texas Press, 1987.
- R. L. Kashyap and M. R. Bell, "Error correcting code-like chanting procedures in ancient India," in *Scientific Heritage of India*, Mythic Society, 1998, pp. 16–29.
- J. Sima and J. Bruck, "Optimal k -deletion correcting codes," *IEEE Trans. Inf. Theory*, vol. 66, pp. 3364–3375, 2020.
- J. Sima and J. Bruck, "On the trace reconstruction problem with a small number of traces," *IEEE Trans. Inf. Theory*, vol. 67, pp. 3529–3544, 2021.
- G. M. Church et al., "Next-generation digital information storage in DNA," *Science*, vol. 337, p. 1628, 2012.
- R. Heckel et al., "A characterization of the DNA data storage channel," *Sci. Rep.*, vol. 9, p. 9663, 2019.
- M. Shirvanimoghaddam et al., "Short code design for URLLC: Challenges and opportunities," *IEEE Commun. Mag.*, vol. 57, pp. 60–66, 2019.
- O. Khatib, "Real-time obstacle avoidance for manipulators and mobile robots," *Int. J. Robot. Res.*, vol. 5, pp. 90–98, 1986.